

```
In [2]: import numpy as np
import pandas as pd
from collections import Counter
import warnings
warnings.filterwarnings('ignore')
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [3]: df=pd.read_csv("diabetes.csv")
```

```
In [4]: df
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.288
...
763	10	101	76	48	180	32.9	0.171
764	2	122	70	27	0	36.8	0.340
765	5	121	72	23	112	26.2	0.245
766	1	126	60	0	0	30.1	0.349
767	1	93	70	31	0	30.4	0.315

768 rows × 9 columns



```
In [5]: df.isnull().sum()
```

```
Out[5]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [6]: df.columns
```

```
Out[6]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [7]: df.head(10)
```

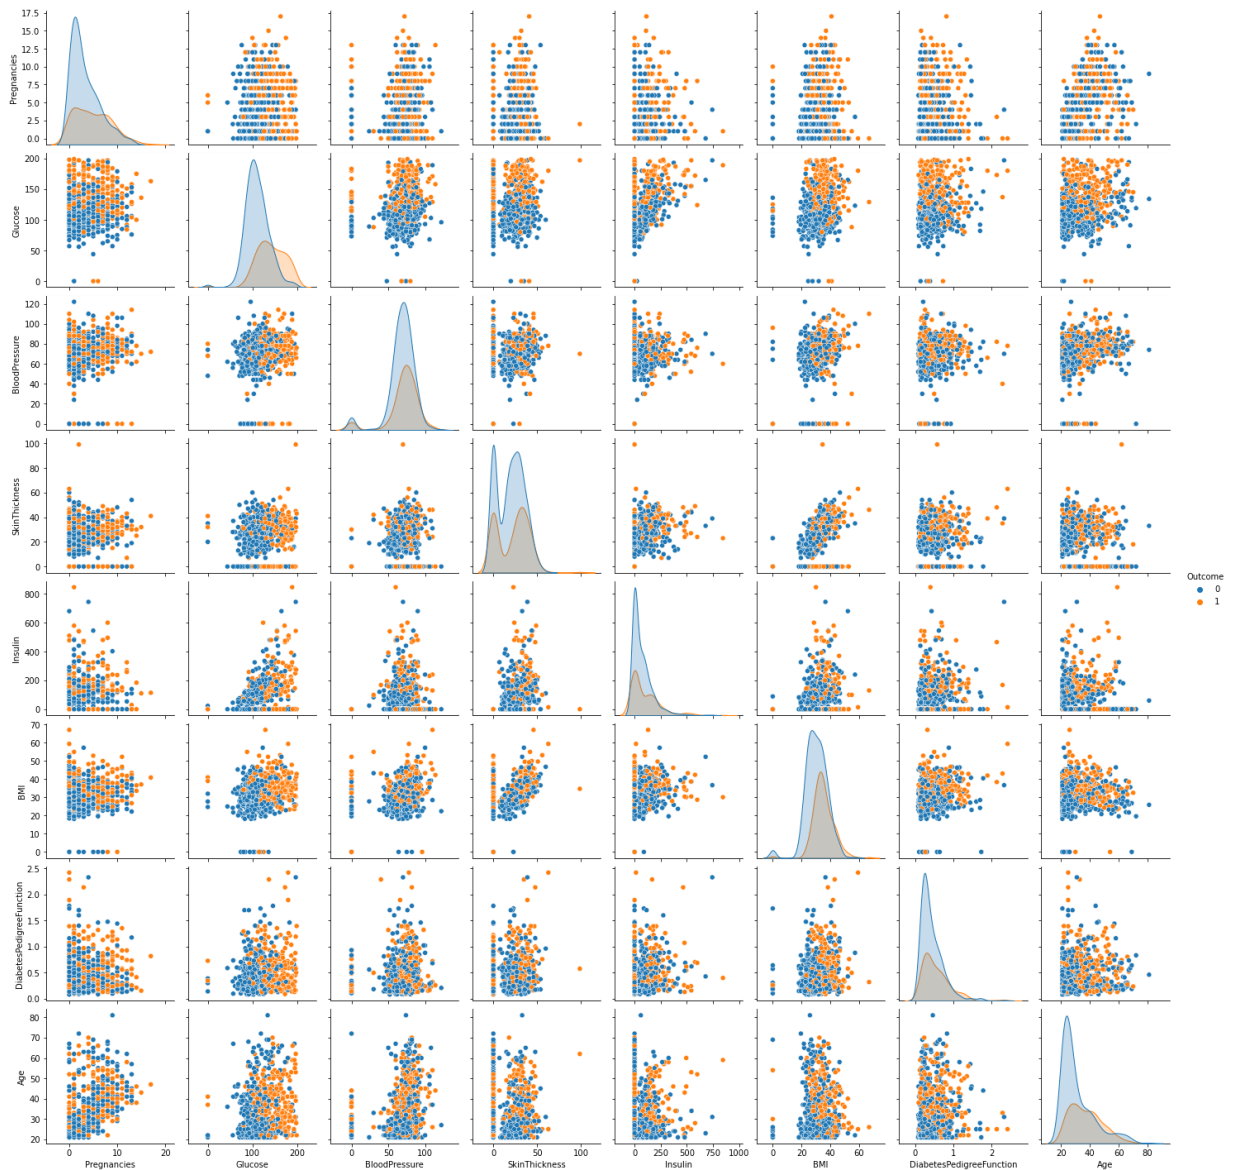
```
Out[7]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	51
1	1	85	66	29	0	26.6	0.351	33
2	8	183	64	0	0	23.3	0.672	33
3	1	89	66	23	94	28.1	0.167	34
4	0	137	40	35	168	43.1	2.288	33
5	5	116	74	0	0	25.6	0.201	30
6	3	78	50	32	88	31.0	0.248	26
7	10	115	0	0	0	35.3	0.134	28
8	2	197	70	45	543	30.5	0.158	26
9	8	125	96	0	0	0.0	0.232	27

Data Visualization

```
In [8]: sns.pairplot(df,hue="Outcome")
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1b29630f4f0>
```



In [9]: df.describe().T

Out[9]:

	count	mean	std	min	25%	50%	75%	
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	1
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	19
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	12
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	9
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	84
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	6
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	8
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	

In [10]: df.describe()

Out[10]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
In [11]: df.corr()
```


Out[11]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140625	0.036228	0.292656
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140625	1.000000	0.461501	0.252474
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036228	0.461501	1.000000	0.733294
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292656	0.252474	0.733294	1.000000




```
In [12]: sns.heatmap(df.corr())
```

Out[12]: <AxesSubplot:>




```
In [13]: sns.scatterplot(y='Age',x='SkinThickness',data=df)
plt.show()
```




```
In [14]: sns.boxplot(x='Age',data=df)
```

```
Out[14]: <AxesSubplot:xlabel='Age'>
```

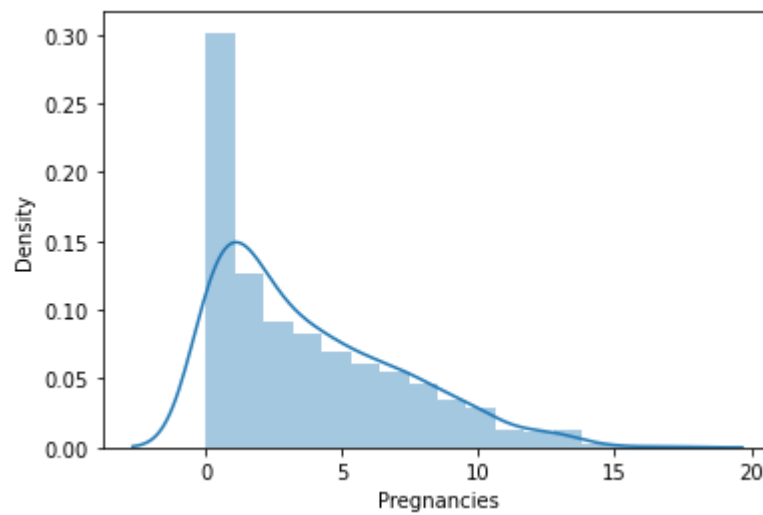


```
In [15]: sns.scatterplot(x='BMI',y='Age',data=df)
plt.show()
```




```
In [16]: sns.distplot(df['Pregnancies'])
```

```
Out[16]: <AxesSubplot:xlabel='Pregnancies', ylabel='Density'>
```




```
In [17]: sns.boxplot(x='Glucose',data=df)
```

```
Out[17]: <AxesSubplot:xlabel='Glucose'>
```





```
In [18]: sns.jointplot(x='BMI',y='Pregnancies',data=df)
```

```
Out[18]: <seaborn.axisgrid.JointGrid at 0x1b299bf68b0>
```




```
In [19]: sns.boxplot(x='Insulin',data=df)
```

```
Out[19]: <AxesSubplot:xlabel='Insulin'>
```




```
In [20]: sns.kdeplot(x='Age',data=df)
```

```
Out[20]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```




```
In [21]: sns.kdeplot(x='Glucose',data=df)
```

```
Out[21]: <AxesSubplot:xlabel='Glucose', ylabel='Density'>
```



```
In [22]: sns.scatterplot(y='SkinThickness',x='Age',data=df)
plt.show()
```



Data Cleaning

```
In [23]: df.describe()
```

```
Out[23]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In the above table there is 0 value in the specific columns that doesn't make any sense, so we have to replace all the zeroes with a nan value for easy reading and understanding, and we can replace this Nan value with some value for better data processing.

There is 0 value in the dataset in columns 1)Glucose 2)BloodPressure 3)SkinThickness 4)Insulin

5) BMI

```
In [24]: df_copy=df.copy(deep=True)
```

```
In [25]: df_copy[['Glucose','BloodPressure','SkinThickness','Insulin','BMI']]=df_copy[['Glucose','BloodPressure','SkinThickness','Insulin','BMI']]
```

```
In [26]: df_copy
```

```
Out[26]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148.0	72.0	35.0	NaN	33.6	0.627
1	1	85.0	66.0	29.0	NaN	26.6	0.351
2	8	183.0	64.0	NaN	NaN	23.3	0.672
3	1	89.0	66.0	23.0	94.0	28.1	0.167
4	0	137.0	40.0	35.0	168.0	43.1	2.288
...
763	10	101.0	76.0	48.0	180.0	32.9	0.171
764	2	122.0	70.0	27.0	NaN	36.8	0.340
765	5	121.0	72.0	23.0	112.0	26.2	0.245
766	1	126.0	60.0	NaN	NaN	30.1	0.349
767	1	93.0	70.0	31.0	NaN	30.4	0.315

768 rows × 9 columns

```
In [27]: df_copy.isnull().sum()
```

```
Out[27]: Pregnancies      0
Glucose      5
BloodPressure 35
SkinThickness 227
Insulin      374
BMI          11
DiabetesPedigreeFunction 0
Age          0
Outcome      0
dtype: int64
```

```
In [28]: df_copy.describe()
```

```
Out[28]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	763.000000	733.000000	541.000000	394.000000	757.000000	
mean	3.845052	121.686763	72.405184	29.153420	155.548223	32.457464	
std	3.369578	30.535641	12.382158	10.476982	118.775855	6.924988	
min	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	
25%	1.000000	99.000000	64.000000	22.000000	76.250000	27.500000	
50%	3.000000	117.000000	72.000000	29.000000	125.000000	32.300000	
75%	6.000000	141.000000	80.000000	36.000000	190.000000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
In [29]: df_copy.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   Pregnancies                          768 non-null    int64  
1   Glucose                              763 non-null    float64  
2   BloodPressure                        733 non-null    float64  
3   SkinThickness                        541 non-null    float64  
4   Insulin                              394 non-null    float64  
5   BMI                                  757 non-null    float64  
6   DiabetesPedigreeFunction             768 non-null    float64  
7   Age                                  768 non-null    int64  
8   Outcome                              768 non-null    int64  
dtypes: float64(6), int64(3)  
memory usage: 54.1 KB
```

```
In [32]: df_copy['Glucose'].fillna(df_copy['Glucose'].mean(),inplace=True)
```

```
In [33]: df_copy.isnull().sum()
```

```
Out[33]: Pregnancies      0  
Glucose      0  
BloodPressure    35  
SkinThickness   227  
Insulin         374  
BMI            11  
DiabetesPedigreeFunction  0  
Age            0  
Outcome        0  
dtype: int64
```

```
In [34]: df_copy['Pregnancies'].fillna(df_copy['Pregnancies'].mean(),inplace=True)
```

```
In [42]: df_copy.isnull().sum()
```

```
Out[42]: Pregnancies      0
         Glucose          0
         BloodPressure    0
         SkinThickness    0
         Insulin          0
         BMI              0
         DiabetesPedigreeFunction  0
         Age              0
         Outcome          0
         dtype: int64
```

```
In [43]: df_copy['BloodPressure'].fillna(df_copy['BloodPressure'].mean(),inplace=True)
```

```
In [44]: df_copy['SkinThickness'].fillna(df_copy['SkinThickness'].median(),inplace=True)
```

```
In [45]: df_copy['Insulin'].fillna(df_copy['Insulin'].median(),inplace=True)
```


```
In [48]: df_copy['BMI'].fillna(df_copy['BMI'].median(),inplace=True)
```

```
In [49]: df_copy.isnull().sum()
```

```
Out[49]: Pregnancies      0
         Glucose          0
         BloodPressure    0
         SkinThickness    0
         Insulin          0
         BMI              0
         DiabetesPedigreeFunction  0
         Age              0
         Outcome          0
         dtype: int64
```

```
In [51]: sns.boxplot(x='BMI',data=df)
```

```
Out[51]: <AxesSubplot:xlabel='BMI'>
```



```
In [53]: df_copy.describe()
```

```
Out[53]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	121.686763	72.405184	29.153420	140.671875	32.457464	
std	3.369578	30.435949	12.096346	8.790942	86.383060	6.875151	
min	0.000000	44.000000	24.000000	7.000000	14.000000	18.200000	
25%	1.000000	99.750000	64.000000	25.000000	121.500000	27.500000	
50%	3.000000	117.000000	72.202592	29.153420	125.000000	32.400000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

Scaling The Data

#WE use the Scaling for dealing with the outlier outlier basically those data which present in outside the graph the outlier will be affected to the our data so that we want standrize our data


```
In [57]: sns.boxplot(x='Outcome',y='BMI',data=df)
```

```
Out[57]: <AxesSubplot:xlabel='Outcome', ylabel='BMI'>
```



```
In [58]: sns.boxplot(y='Insulin',x='Outcome',data=df)
```

```
Out[58]: <AxesSubplot:xlabel='Outcome', ylabel='Insulin'>
```



```
In [62]: from sklearn.preprocessing import StandardScaler  
scaler=StandardScaler()
```



```
In [63]: scaler.fit(df_copy.drop('Outcome',axis=1))
```

```
Out[63]: StandardScaler()
```

```
In [69]: scaled_features=scaler.transform(df.drop('Outcome',axis=1))
```

```
In [70]: scaled_features
```

```
Out[70]: array([[ 0.63994726,  0.86510807, -0.03351824, ...,  0.16629174,
                  0.46849198,  1.4259954 ],
                [-0.84488505, -1.20616153, -0.52985903, ..., -0.85253118,
                  -0.36506078, -0.19067191],
                [ 1.23388019,  2.0158134 , -0.69530596, ..., -1.33283341,
                  0.60439732, -0.10558415],
                ...,
                [ 0.3429808 , -0.0225789 , -0.03351824, ..., -0.91074963,
                  -0.68519336, -0.27575966],
                [-0.84488505,  0.14180757, -1.02619983, ..., -0.34311972,
                  -0.37110101,  1.17073215],
                [-0.84488505, -0.94314317, -0.19896517, ..., -0.29945588,
                  -0.47378505, -0.87137393]])
```

```
In [73]: df_stand=pd.DataFrame(scaled_features,columns=df.columns[:-1])
df_stand.head(10)
```

```
Out[73]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	0.639947	0.865108	-0.033518	0.665502	-1.629527	0.166292	0.4
1	-0.844885	-1.206162	-0.529859	-0.017463	-1.629527	-0.852531	-0.5
2	1.233880	2.015813	-0.695306	-3.318463	-1.629527	-1.332833	0.6
3	-0.844885	-1.074652	-0.529859	-0.700429	-0.540642	-0.634212	-0.9
4	-1.141852	0.503458	-2.680669	0.665502	0.316566	1.548980	5.4
5	0.342981	-0.186965	0.131929	-3.318463	-1.629527	-0.998077	-0.8
6	-0.250952	-1.436303	-1.853434	0.324019	-0.610145	-0.212128	-0.6
7	1.827813	-0.219843	-5.989608	-3.318463	-1.629527	0.413720	-1.0
8	-0.547919	2.476096	-0.198965	1.803778	4.660524	-0.284901	-0.9
9	1.233880	0.108930	1.951845	-3.318463	-1.629527	-4.724058	-0.7

```
In [74]: df_stand['Outcome']=df['Outcome']
```

```
In [75]: df_stand.head(10)
```


```
Out[75]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	0.639947	0.865108	-0.033518	0.665502	-1.629527	0.166292	0.4
1	-0.844885	-1.206162	-0.529859	-0.017463	-1.629527	-0.852531	-0.5
2	1.233880	2.015813	-0.695306	-3.318463	-1.629527	-1.332833	0.6
3	-0.844885	-1.074652	-0.529859	-0.700429	-0.540642	-0.634212	-0.9
4	-1.141852	0.503458	-2.680669	0.665502	0.316566	1.548980	5.4
5	0.342981	-0.186965	0.131929	-3.318463	-1.629527	-0.998077	-0.8
6	-0.250952	-1.436303	-1.853434	0.324019	-0.610145	-0.212128	-0.6
7	1.827813	-0.219843	-5.989608	-3.318463	-1.629527	0.413720	-1.0
8	-0.547919	2.476096	-0.198965	1.803778	4.660524	-0.284901	-0.9
9	1.233880	0.108930	1.951845	-3.318463	-1.629527	-4.724058	-0.7




```
In [76]: sns.boxplot(x='Outcome',y='BMI',data=df_stand)
```

```
Out[76]: <AxesSubplot:xlabel='Outcome', ylabel='BMI'>
```



```
In [77]: sns.boxplot(x='Outcome',y='Insulin',data=df_stand)
```

```
Out[77]: <AxesSubplot:xlabel='Outcome', ylabel='Insulin'>
```



Dependent And Independent Set

```
In [81]: x=df_stand.drop('Outcome',axis=1)
```

In [82]:

x

Out[82]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	0.639947	0.865108	-0.033518	0.665502	-1.629527	0.166292	
1	-0.844885	-1.206162	-0.529859	-0.017463	-1.629527	-0.852531	-
2	1.233880	2.015813	-0.695306	-3.318463	-1.629527	-1.332833	
3	-0.844885	-1.074652	-0.529859	-0.700429	-0.540642	-0.634212	-
4	-1.141852	0.503458	-2.680669	0.665502	0.316566	1.548980	
...	
763	1.827813	-0.680125	0.297376	2.145261	0.455573	0.064409	-
764	-0.547919	0.010298	-0.198965	-0.245119	-1.629527	0.632039	-
765	0.342981	-0.022579	-0.033518	-0.700429	-0.332132	-0.910750	-
766	-0.844885	0.141808	-1.026200	-3.318463	-1.629527	-0.343120	-
767	-0.844885	-0.943143	-0.198965	0.210192	-1.629527	-0.299456	-

768 rows × 8 columns



In [83]:

y=df_stand['Outcome']

In [84]:

y

Out[84]:

```
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
```

Name: Outcome, Length: 768, dtype: int64

Train Test Split

In [86]:

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [87]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 537 entries, 537 to 648
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          537 non-null    float64
1   Glucose                              537 non-null    float64
2   BloodPressure                        537 non-null    float64
3   SkinThickness                       537 non-null    float64
4   Insulin                             537 non-null    float64
5   BMI                                 537 non-null    float64
6   DiabetesPedigreeFunction             537 non-null    float64
7   Age                                 537 non-null    float64
dtypes: float64(8)
memory usage: 37.8 KB
```

```
In [90]: X_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 231 entries, 152 to 349
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          231 non-null    float64
1   Glucose                              231 non-null    float64
2   BloodPressure                        231 non-null    float64
3   SkinThickness                       231 non-null    float64
4   Insulin                             231 non-null    float64
5   BMI                                 231 non-null    float64
6   DiabetesPedigreeFunction             231 non-null    float64
7   Age                                 231 non-null    float64
dtypes: float64(8)
memory usage: 16.2 KB
```

Logistic Regression Algorithms

```
In [92]: from sklearn.linear_model import LogisticRegression
```

```
In [93]: logmodel=LogisticRegression()
```

```
In [95]: logmodel.fit(X_train,y_train)
```

```
Out[95]: LogisticRegression()
```

```
In [96]: predictions=logmodel.predict(X_test)
```

```
In [97]: from sklearn.metrics import classification_report
```

```
In [99]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.79	0.91	0.85	152
1	0.75	0.54	0.63	79
accuracy			0.78	231
macro avg	0.77	0.73	0.74	231
weighted avg	0.78	0.78	0.77	231

```
In [100]: from sklearn.metrics import confusion_matrix
```

```
In [101]: print(confusion_matrix(y_test,predictions))
```

```
[[138  14]
 [ 36  43]]
```

```
In [ ]:
```